

# Josh Goldstein

Principal Solutions Architect | AI Infrastructure & Solutions Leader

76 Roebling St, Apt 1B, Brooklyn, NY 11211 305-321-6144 Josh4goldstein@gmail.com

<https://linkedin.com/in/joshsgoldstein> <https://github.com/joshsgoldstein>

## HIGHLIGHTS

Technical leader with over 10 years of experience architecting and scaling AI infrastructure and enterprise search solutions from pre-sales through production. Proven track record of building large-scale data solutions, influencing product roadmaps, and supporting multi-million dollar opportunities across Fortune 500 companies and high-growth startups. Expertise in information retrieval and LLM/MLOps, with a passion for translating complex technology into business outcomes. Skilled at bridging engineering, product, and GTM teams to accelerate customer success and revenue growth.

## WORK HISTORY

**Customer Excellence Solutions Engineer** Aug 2024 – Oct 2025 · Remote / Amsterdam, ND  
— Weaviate

- Defined and launched the Bring Your Own Cloud (BYOC) offering, closing a Fortune 50 enterprise enabling a centralized platform for GenAI applications and defined the enterprise product serving multiple large enterprise customers.
- Partnered cross-functionally to strengthen enterprise positioning through customer feedback, technical enablement content (Jupyter notebooks, tutorials, eBooks), and conference presentations, while advising teams and clients on emerging AI capabilities.
- Advised and architected complex agentic and semantic search solutions integrating multimodal retrieval and generative pipelines, showing what's possible and bringing enterprise use cases from concept into production enabling successful product launches and updates to legacy enterprise systems.
- Mentored new Solutions Engineers and established standardized POC framework and product feedback processes to accelerate enterprise delivery and adoption.

**Senior Solutions Architect — Seldon** Jan 2022 – Jul 2024 · Remote / London, UK

- Established the North America Solutions Engineering function, creating repeatable customer engagement frameworks, POC methodologies, and post-sales delivery patterns that improved time-to-value across enterprise and startup customers.
- Co-designed Seldon's LLM product offering, defining agentic AI patterns, customer experience flows, integration models, and establishing best practice patterns to enable the adoption of generative AI.
- Contributed directly to MLServer (core product) by fixing production issues for customers leading to a \$500k+ renewal of an at-risk enterprise account.
- Designed and implemented enterprise-grade inference systems and deployment patterns for ML Endpoints, improving reliability, scalability, and time to production for customers deploying machine learning models at scale and deploying GenAI applications.
- Represented customer needs in engineering discussions and closed feedback loops that strengthened Seldon's presence and adoption in the MLOps community.

**Senior Solutions Architect — Keeeb** Jan 2021 – Dec 2021 · Remote / Boston, MA

- Closed the company's first U.S. enterprise deal in the manufacturing sector by designing and delivering a proof of concept that integrated internal data sources into the company's internal enterprise content discovery platform.
- Led technical and go-to-market execution for Keeeb's federated search platform, overseeing pre-sales, customer enablement, and solution design for enterprise accounts.
- Contributed to product connectors that improved data accessibility, demonstrated agile engineering execution, and unlocked key features required to close first U.S. enterprise deal.
- Structured and led a joint engagement with a systems integrator to migrate a customer's legacy data system to the Microsoft ecosystem, aligning technical

## CORE COMPETENCIES

Customer Success & Solution Delivery, Cross-Functional Leadership, Product & GTM Strategy Alignment, Technical Evangelism & Public Speaking, Team Mentorship & Enablement, Requirements Gathering & Solution Design, Strategic Storytelling & Communication, Self-Motivated & Outcome-Driven

## TECHNICAL SKILLS

AI Infrastructure & Enterprise Architecture, Agentic Application Design (RAG frameworks, context management), Vector Databases & Semantic Search (Multimodal, Hybrid Search), LLM & Generative AI Deployment, MLOps Platform Architecture, Kubernetes & Cloud-Native Systems (AWS / GCP / Azure), CI/CD & Observability, Python Application Development, Data Engineering & Pipeline Design, Authentication & Authorization, Database Technologies (SQL / NoSQL / Graph), Product Feedback Loops & Experimentation

## SPEAKING

**Purpose-Built Data Agents: This Is the Way**  
MLOps World / Gen AI Summit - October 2025

**The Value of Data Agentic Tooling: A Work in Progress**  
MLOps World / Gen AI Summit - August 2025

**WIP: Data and the Value of Data Agentic Tooling**  
Toronto Machine Learning Summit - June 2025

**Digital Experience for the Next Generation Knowledge Worker**  
Activate 2019

**Common Mistakes Made Deploying a Search Application with Lucidworks Fusion**  
Activate 2018

**Ethics in Tech**  
UC Berkeley EECS Honors Program

**From Model to Microservice > MLOps at Scale**  
Southern Data Science Conference

**Various Apache Solr Meetups**

## EDUCATION

**Bachelor of Science: Information Technology**

scope, pricing, and delivery to ensure success.

**Business Innovation Manager —  
Lucidworks Inc.**

Mar 2019 – Dec 2019 · San  
Francisco, CA

- Partnered with Product and Engineering leadership to design and prototype new capabilities for Lucidworks Fusion, drawing on experience as one of the platform's first enterprise customers and an early subject-matter expert.
- Helped define and deliver key product initiatives including a product merchandiser for digital commerce and a low-code/no-code framework for building search applications, contributing to both the initial and next-generation versions.
- Led Lucidworks' first beta program for new products, coordinating with early adopters to capture feedback that informed roadmap priorities and drove product-market fit.
- Collaborated with Customer Success and field teams to turn customer insights into product requirements, driving retention and expansion within enterprise accounts.

**Solution Architect — Lucidworks Inc. Jun 2017 – Mar 2019 · San Francisco, CA**

- Designed and implemented enterprise search and knowledge solutions using Lucidworks Fusion and Apache Solr for Fortune 500 customers across technology and energy sectors.
- Influenced product direction by relaying customer needs to Product teams and validating new Fusion capabilities with key accounts, ensuring field feedback shaped development priorities.
- Integrated third-party technologies and vendor tools into the Lucidworks stack to expand solution capabilities and support complex customer environments enabling a fortune 10 energy company to unlock enterprise search for the entire organization .
- Created product demos for sales and marketing that contributed to Lucidworks' inclusion in the Gartner Magic Quadrant for Insight Engines.
- Mentored new Solution Architects following the Twigkit acquisition, training them in search theory, Fusion, and Solr technologies.
- Developed internal enablement programs and workshops to strengthen technical expertise across customer-facing teams and introduce new features in the Lucidworks product suite.

**Sales Engineer — Lucidworks Inc. Feb 2016 – Jun 2017 · San Francisco, CA**

- Supported sales teams by designing and presenting search-based solutions to enterprise clients, leading proof-of-concept projects for opportunities ranging from \$100K-\$1M+.
- Directly contributed to 15 consecutive quarters of sales growth by delivering technical wins in complex enterprise evaluations.
- Acted as a technical liaison between customers and engineering, ensuring field feedback informed early product releases and feature prioritization.
- Enhanced Lucidworks' training program by leveraging AWS to improve reliability, scalability, and instructional content.

**Co-Founder / Technical Lead — Mulsanne  
Dashboards**

Apr 2014 – Jan 2016 ·  
Houston, TX

- Designed and led development of Mulsanne Dashboards, an executive-level news and analytics platform that delivered personalized, business-critical updates to help leaders make better decisions throughout the day.
- Architected and implemented the platform's end-to-end infrastructure using Apache Solr, ExpressJS, and MongoDB, supporting real-time data ingestion, indexing, and search.
- Built an NLP/ML data pipeline that aggregated and classified news using RSS feeds, web scraping, NLP techniques, and bespoke machine-learning models developed with an external partner.
- Created internal monitoring and alerting dashboards to ensure system reliability and data accuracy.
- Defined product design, user workflows, and go-to-market strategy in collaboration with the founding team.

Rochester Institute of Technology,  
Rochester, NY  
Aug 2007 – Dec 2011  
Concentrations in Finance and  
Communications.